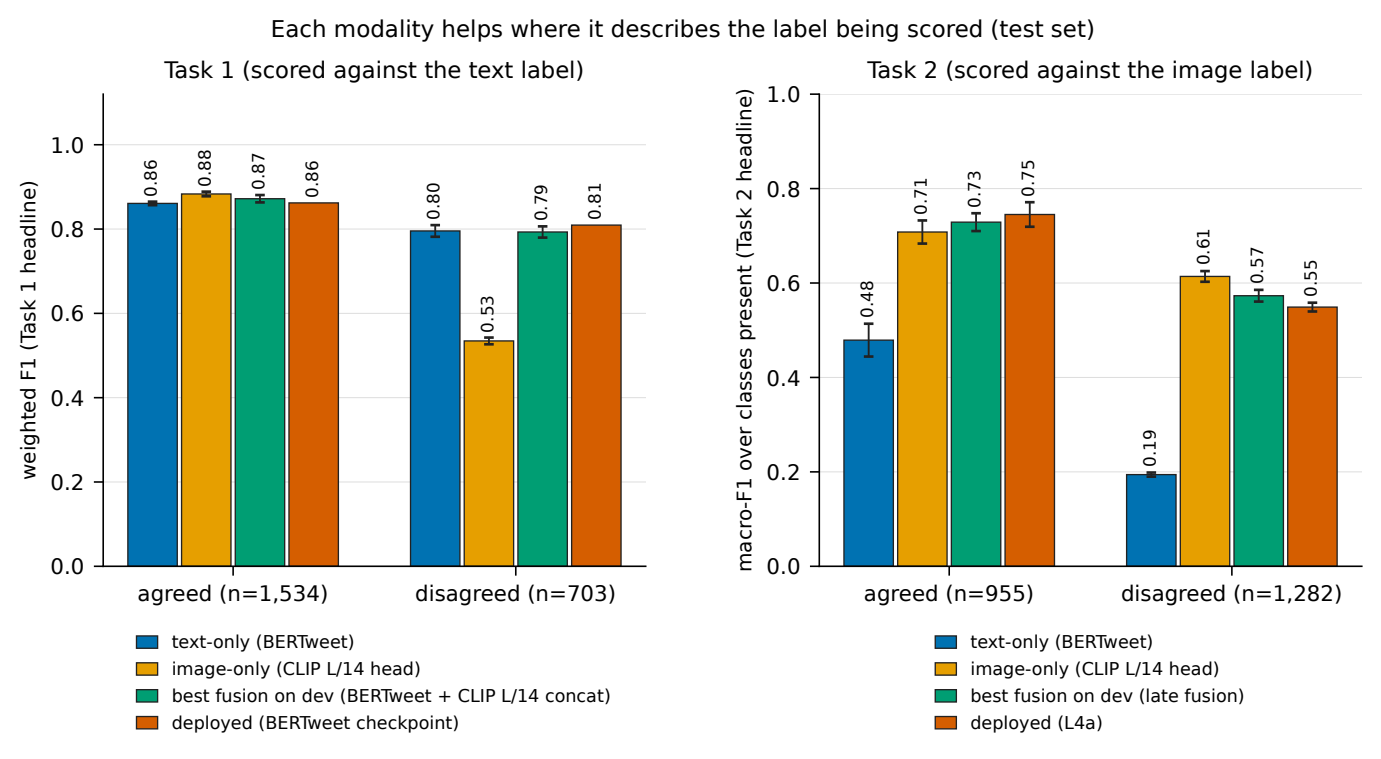
Each modality helps where it describes the label being scored (test set)
Task 1 (scored against the text label)
weighted F1 (Task 1 headline)
1.0
0.8
0.6
0.4
0.2
0.0
0.86
0.88
0.87
0.86
0.80
0.53
0.79
0.81
agreed (n=1,534)
disagreed (n=703)
text-only (BERTweet)
image-only (CLIP L/14 head)
best fusion on dev (BERTweet + CLIP L/14 concat)
deployed (BERTweet checkpoint)
Task 2 (scored against the image label)
macro-F1 over classes present (Task 2 headline)
1.0
0.8
0.6
0.4
0.2
0.0
0.48
0.71
0.73
0.75
0.19
0.61
0.57
0.55
agreed (n=955)
disagreed (n=1,282)
text-only (BERTweet)
image-only (CLIP L/14 head)
best fusion on dev (late fusion)
deployed (L4a)