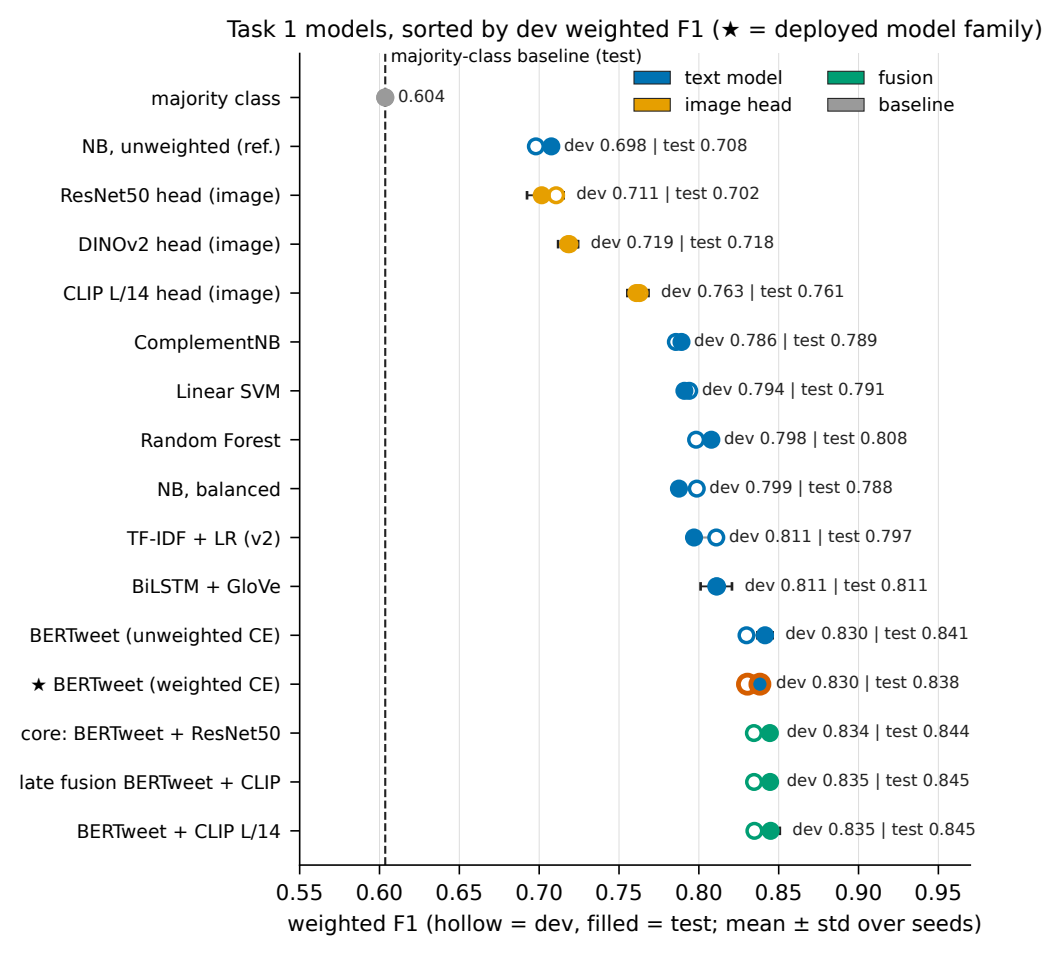
Task 1 models, sorted by dev weighted F1 (★ = deployed model family)
majority-class baseline (test)
text model
image head
fusion
baseline
majority class
0.604
NB, unweighted (ref.)
dev 0.698 | test 0.708
ResNet50 head (image)
dev 0.711 | test 0.702
DINOv2 head (image)
dev 0.719 | test 0.718
CLIP L/14 head (image)
dev 0.763 | test 0.761
ComplementNB
dev 0.786 | test 0.789
Linear SVM
dev 0.794 | test 0.791
Random Forest
dev 0.798 | test 0.808
NB, balanced
dev 0.799 | test 0.788
TF-IDF + LR (v2)
dev 0.811 | test 0.797
BiLSTM + GloVe
dev 0.811 | test 0.811
BERTweet (unweighted CE)
dev 0.830 | test 0.841
★ BERTweet (weighted CE)
dev 0.830 | test 0.838
core: BERTweet + ResNet50
dev 0.834 | test 0.844
late fusion BERTweet + CLIP
dev 0.835 | test 0.845
BERTweet + CLIP L/14
dev 0.835 | test 0.845
0.55
0.60
0.65
0.70
0.75
0.80
0.85
0.90
0.95
weighted F1 (hollow = dev, filled = test; mean ± std over seeds)