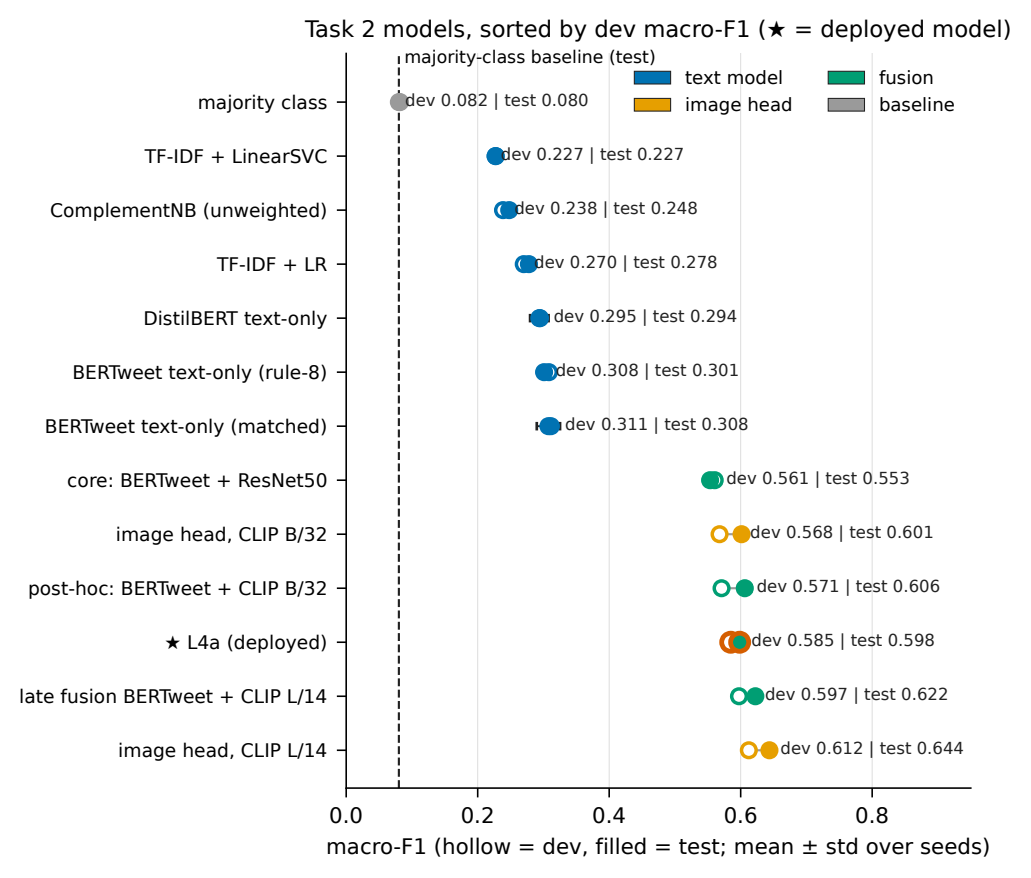
Task 2 models, sorted by dev macro-F1 (★ = deployed model)
majority-class baseline (test)
text model
image head
fusion
baseline
majority class
dev 0.082 | test 0.080
TF-IDF + LinearSVC
dev 0.227 | test 0.227
ComplementNB (unweighted)
dev 0.238 | test 0.248
TF-IDF + LR
dev 0.270 | test 0.278
DistilBERT text-only
dev 0.295 | test 0.294
BERTweet text-only (rule-8)
dev 0.308 | test 0.301
BERTweet text-only (matched)
dev 0.311 | test 0.308
core: BERTweet + ResNet50
dev 0.561 | test 0.553
image head, CLIP B/32
dev 0.568 | test 0.601
post-hoc: BERTweet + CLIP B/32
dev 0.571 | test 0.606
★ L4a (deployed)
dev 0.585 | test 0.598
late fusion BERTweet + CLIP L/14
dev 0.597 | test 0.622
image head, CLIP L/14
dev 0.612 | test 0.644
0.0
0.2
0.4
0.6
0.8
macro-F1 (hollow = dev, filled = test; mean ± std over seeds)